

YOUR CODING AGENT SAYS

“The prompt is broken.”

An AI failure can come from **9** places.

Only **1** of them is the prompt.



Titration: let your agent fix a prompt until it actually works, and prove it did.

MCP SERVER

18 TOOLS

open source · self-hosted

Claude Code · Codex · Cursor

SWIPE →

THE NINE FAILURE ORIGINS

The periodic table of *agent failure.*

01 Su system-under-test fix → the prompt	02 Cg corpus-gap fix → the tests	03 Ra rubric-ambiguity fix → the rubric
04 Jv judge-variance fix → the judges	05 Fm formatter fix → the code	06 If instrument-failure fix → the harness
07 Pp pre-existing not this change	08 Sa state-artifact fix → upstream	09 Ce cost-exhausted fix → the budget

Only **01** means “edit the prompt.”

classify_failure

SWIPE →

HOW IT DECIDES

Three vendors judge. A split gets *argued*.

Judge A

vendor 1

Fm

Judge B

vendor 2

Fm

Judge C

vendor 3

Su



RECONSIDER · 2-1 SPLIT

A and B must answer C's dissent:

STICKING no real defect found



Fm

Formatter.

→ fix the code, leave the prompt alone

your agent's vendor never judges

SWIPE →

THE IMPROVEMENT LOOP

One change at a time,
against a baseline that
doesn't move.

REAL RUN • EXAMPLES/TICKET-TRIAGE

90% → 0%

Billing tickets wrongly marked urgent, fixed in one measured change. 90 of 90 judge votes.

i. The model never grades its own work.

ii. No agreement, no score.

iii. Every lesson becomes memory.

OPEN SOURCE • APACHE-2.0

Runs on *your machine.*

```
$ git clone ../titration
$ docker compose up -d
$ npm install
$ npm run setup
```

github.com/kaitthoughtarchitect/titration

Star it, fork it, or start with the worked example in **examples/ticket-triage.**

“...you must not fool yourself and you are the easiest person to fool.”

RICHARD FEYNMAN

self-hosted • any MCP client